



ANSIBLE DAILY TASKS

ANSIBLE PLAYBOOKS & MODULES

- COMPERHENSIVE GUIDE
- THEORY + PRACITCAL TASKS
- REAL TIME SCENARIO TASKS
- END TO END DOC

Contact Us:

fenilgajjar.devops@gmail.com linkedin.com/in/fenil-gajjar github.com/Fenil-Gajjar



The DevOps Guide to

Ansible Playbooks &





Welcome & Thank You!

First of all — **a huge thank you** to everyone who has followed and supported my previous documentation series A

Your feedback, encouragement, and enthusiasm keep this journey alive. I truly appreciate the time you've spent reading, sharing, and learning from the content I've created.

Now, I'm excited to bring you another value-packed resource:

Automate Like a Pro: The DevOps Guide to Ansible Playbooks & Modules

This guide is more than just a walkthrough — it's a **complete hands-on journey** through the core of what makes Ansible powerful:

Playbooks and Modules — the very heart of infrastructure automation.

You'll learn not just what they are, but:

- Mean of the properties
 Mean of the prope
- Q How to structure them properly
- K How to use them in real-world tasks DevOps engineers actually face
- Plus: best practices, troubleshooting tips, and reusable automation patterns

Whether you're just getting started or looking to sharpen your skills, this guide is designed to be **clear**, **practical**, **and production-minded** — exactly how modern DevOps documentation should be.

Let's get into it — and as always, thank you again for reading! 🙌

Introduction to Ansible Playbooks

Ansible is a powerful automation tool used for configuration management, application deployment, and orchestration. At the heart of this automation lies the **Ansible Playbook**—a YAML-based file that defines **what** needs to be done and **how** to do it, in a clear and structured manner.

What is a Playbook?

An **Ansible Playbook** is a blueprint of automation tasks. It describes a list of actions (called "plays") that need to be executed on one or more remote systems. Each play targets a group of hosts and defines a set of **tasks**, which are executed using **Ansible modules**.

Think of playbooks as human-readable instructions for machines—concise, repeatable, and scalable.

Why Use Playbooks?

Here's why playbooks are essential in modern infrastructure management:

- Consistency: Same set of tasks applied across environments without manual intervention.
- Idempotency: Re-running playbooks doesn't break the system; it ensures the desired state.

- Clarity: Written in YAML, playbooks are easy to read, review, and version-control.
- Declarative Style: You describe the end state, and Ansible figures out how to get there.
- **Extensibility**: Easily integrate with roles, variables, templates, and more for larger projects.

Mark How Playbooks Work

Playbooks work by:

- 1. Connecting to the defined hosts (via SSH or WinRM).
- 2. Running each task sequentially using appropriate modules.
- 3. Ensuring that each task changes the system only if needed (idempotent behavior).
- 4. Reporting results in a clean and informative way.

A Simple Example

- name: Install and start Apache web server hosts: webservers become: yes tasks: - name: Install Apache apt: name: apache2 state: present - name: Start Apache service service: name: apache2 state: started

enabled: yes

This playbook:

- Runs on hosts in the webservers group.
- Installs Apache using the apt module.
- Ensures the Apache service is running and enabled at boot using the service module.

% When to Use Playbooks

Playbooks are ideal when:

- You need to automate a **sequence** of tasks.
- You want to manage configuration at scale.
- You're implementing **CI/CD** pipelines or **infra-as-code**.
- You want a **version-controlled** approach to infrastructure.

(6) Importance of Playbooks in Automation

Ansible playbooks are the **core engine** behind infrastructure automation. While ad-hoc commands are suitable for one-time tasks, playbooks provide the **repeatability**, **structure**, and **reliability** that enterprise-grade automation demands.

Repeatability & Consistency

Playbooks ensure that automation tasks execute the same way every time, across all environments—development, staging, and production. This eliminates "configuration drift" and manual inconsistencies.

Example: Installing NGINX on 10 servers with the same config? One playbook does it uniformly, every time.

Declarative Infrastructure as Code (IaC)

Playbooks allow you to define **desired state**, not the step-by-step instructions to get there. This abstracts complexity and aligns with modern DevOps practices like IaC.

You declare: "NGINX must be installed and running."

Ansible ensures: That's the final result, regardless of the server's current state.

Scalability Across Environments

With inventory grouping and variables, playbooks can target:

- A single host
- A specific group of servers
- An entire data center or cloud region

This makes it easy to scale operations from local VMs to thousands of machines.

Documentation & Auditing

Playbooks double as **living documentation** of your infrastructure. Anyone on your team can read and understand:

- What configurations are applied
- When changes were made (via version control)
- Why certain tasks exist (with descriptions/comments)



Safe, Idempotent Execution

Playbooks can be safely re-run multiple times without causing adverse effects. Ansible checks whether changes are needed before taking action, preventing disruptions or duplication.

Re-running a playbook won't reinstall a package if it's already there—it only makes changes when necessary.

Integration with DevOps Toolchains

Playbooks can be easily integrated into:

- CI/CD pipelines (e.g., Jenkins, GitLab CI)
- Cloud provisioning tools (e.g., Terraform)
- **Monitoring systems** (e.g., via Ansible callbacks)

This makes automation part of the **entire software lifecycle**, not just system setup.

Fast Rollouts & Disaster Recovery

Need to update 100 servers in minutes? Or reconfigure systems after an outage? Playbooks make it fast, safe, and auditable to push changes or recover systems.



What Are Ansible Modules?

Ansible modules are the **core building blocks** of automation in Ansible. Every task you define in a playbook uses a module under the hood to perform a specific action—like installing packages, managing users, copying files, or restarting services.

You can think of a module as a small, standalone piece of code that knows how to do one thing well.

Key Characteristics of Ansible Modules

- Task-oriented: Each module is designed for a particular job (e.g., install software, create users, manage services).
- **Idempotent**: Modules are built to ensure that running the same task multiple times doesn't cause unintended changes.
- Cross-platform: Many modules work on Linux, Windows, and cloud platforms.
- Reusable: The same module can be used across multiple playbooks and environments.

Types of Modules

Modules are categorized based on their purpose:

Category	Description	Examples
System	Manage OS-level components	user, group, hostname
Package Manager	Install/remove software packages	apt, yum, dnf, pip
File	Work with files and directories	copy, file, template
Service	Start/stop/manage services	service, systemd
Cloud	Manage cloud infrastructure	ec2, azure_rm, gcp_compute_instance
Networking	Configure network devices or settings	<pre>ios_config, nmcli, firewalld</pre>
Utility	Help with debugging, pauses, conditions	debug, pause, assert
Custom	User-created modules for custom use cases	Written in Python or any language

How Modules Are Used in a Playbook

Every **task** in a playbook calls a module. For example:

- name: Install NGINX

apt:

name: nginx

state: present

Here:

- apt is the module.
- It installs the nginx package and ensures it's present.
- Ansible handles the logic behind checking the current state and acting accordingly.

Behind the Scenes

When you run a playbook:

1. Ansible connects to the remote host.

- 2. It sends the module code (or uses the built-in one if local).
- 3. The module executes the task on the target system.
- 4. The output is returned in JSON format to Ansible for processing and display.

Writing Custom Modules

If your organization has unique requirements, you can write your own modules using:

- Python (most common)
- Bash or any executable scripting language
- Return output in JSON format for Ansible to parse

- Modules are the **action layer** of Ansible.
- They abstract away complex operations into simple YAML instructions.
- Understanding which module to use—and how—is key to writing effective automation.

Relationship Between Playbooks and Modules

Understanding the relationship between Ansible Playbooks and Modules is fundamental to mastering Ansible automation. In short:

Playbooks describe what to do, and Modules define how it's done.

Playbooks: The Blueprint

A playbook is a YAML-formatted file that outlines the desired state of your systems. It tells Ansible:

- Which hosts to target
- What tasks to run
- In what **order**
- With what conditions, variables, and privileges

But the playbook itself doesn't do the actual work—it delegates that to modules.

Nodules: The Workers

A **module** is a small program that performs a specific task on the remote system:

Installing a package

- Creating a user
- Modifying a config file
- Restarting a service

Modules are **invoked** by tasks inside the playbook.

X How They Work Together

Each **task** in a playbook uses a module. For example:

- name: Ensure nginx is installed

apt:

name: nginx

state: present

In this example:

- The **playbook** defines a task: "Ensure nginx is installed".
- The **task** uses the apt **module** to perform the actual installation.

 Ansible executes the module code on the remote machine, and reports back success/failure.

Execution Flow Summary

- 1. You write a playbook with tasks.
- 2. **Each task** specifies a module and arguments.
- 3. Ansible connects to the remote host and runs the module.
- 4. The module **performs the action** and returns results to Ansible.
- 5. Ansible **logs and displays** the output to you.

Analogy

Imagine a playbook as a **movie script** and modules as the **actors**:

- The script says what scenes need to happen (tasks).
- The actors (modules) perform the actions.
- The director (Ansible) ensures the performance is in order, repeats if needed, and gives feedback.

Why This Relationship Matters

- It allows separation of logic and execution.
- You can reuse modules across many tasks and playbooks.
- Modules handle complexity; you just declare the outcome in YAML.
- Custom modules extend this relationship for your organization's specific needs.

So, a playbook **orchestrates**, while modules **execute**. Together, they form the complete automation engine of Ansible.

Understanding Ansible Playbooks: A Complete Guide

Ansible playbooks are the **heart** of automation in Ansible. They're how you tell Ansible what you want your systems to look like. In this section, we'll break it all down—from syntax to structure to real-world tips.

YAML Syntax Basics (Start Here)

Playbooks are written in **YAML** (YAML Ain't Markup Language), a simple and human-readable data format. YAML uses **indentation** and **key-value pairs** to describe configurations.

✓ Basic YAML Rules

- Use **spaces**, not tabs. Indentation is crucial.
- Use: to separate keys and values.
- Lists begin with (dash and space).

Example

- name: Simple YAML example

tasks:

- name: Install nginx

apt:

name: nginx

state: present

₹ Tip: If your playbook fails unexpectedly, check your spacing first—it's the #1 cause of syntax errors!

Anatomy of a Playbook

A playbook is made of **plays**, and each play runs a sequence of **tasks** on a group of **hosts**.

Basic Components of a Playbook

Here's what makes up a typical playbook:

_ _ _

- name: Configure web server

hosts: webservers

become: yes # Run with sudo privileges

vars:

http_port: 80

tasks:

- name: Install Apache

apt:

name: apache2

state: present

- name: Ensure Apache is running

service:

name: apache2

state: started

enabled: yes

Breakdown of Key Sections:

name Human-readable name for the play/task

host Target group(s) of machines to run the

Purpose

s play

Section

beco Escalate privileges (e.g., sudo)

me

vars Define custom variables

task List of actions to perform
s

modu Each task uses a module like apt,
les service, etc.

Multi-Play Playbooks

Playbooks can contain **multiple plays** to target different groups of servers in a single run.



- name: Configure web servers

hosts: webservers

become: yes

tasks:

- name: Install nginx

apt:

name: nginx

state: present

- name: Configure database servers

hosts: dbservers

become: yes

tasks:

- name: Install MySQL

apt:

name: mysql-server

state: present

Think of this as automating different roles (web, DB, etc.) in a single deployment plan.

Idempotency in Ansible

One of Ansible's most powerful features is **idempotency**—a fancy word that means:

• "If you run the same playbook 100 times, your system will stay in the same state after the first run unless something actually changes."

Why It Matters

- Avoids duplicate tasks (e.g., doesn't reinstall a package if it's already present).
- Safe to re-run playbooks in CI/CD pipelines or on servers.
- Helps prevent configuration drift over time.

Example

- name: Create a user

user:

name: devops

state: present

This will:

- Add the user if not present.
- Do nothing if the user already exists.
- Never create duplicate users.

Tip: Most core Ansible modules are designed to be idempotent. Custom scripts may not be.

Neal-World Playbook Tips

- Comment generously. Future-you or your teammates will thank you.
- **Use variables** to avoid hardcoding values (http_port, app_version, etc.).
- **Split roles** when your playbooks get big. Use roles/ to organize them.
- Always test your playbook with --check and --diff when possible.

Understanding Ansible Modules: A Complete Guide

Ansible modules are the **workhorses** of your playbooks. While the playbook defines *what* you want, modules define *how it gets done*. Whether you're installing software, copying files, managing services, or configuring cloud resources—a **module is involved in every task**.

What is a Module in Ansible?

A **module** in Ansible is a small, reusable unit of code used to perform a specific task on a remote system. Every **task** you write in a playbook uses one module.

Think of a module as a function that Ansible calls to do something on your target machine.

Example:

- name: Install nginx

apt:

name: nginx

state: present

In this task:

- apt is the module.
- name and state are arguments passed to that module.

Ansible sends the task to the target host, runs the module there, and gets back a result (success/failure/changes made).

T Core Modules vs. Custom Modules

Core Modules

These are the **built-in**, officially supported modules that come with Ansible by default. They are well-tested and regularly maintained.

• Examples: copy, file, apt, yum, service, user, command

Custom Modules

You can create your own modules if:

- A built-in module doesn't meet your need.
- You have a highly specific task or integration.
- You want to wrap custom logic in a reusable form.

Custom modules can be written in:

- Python (recommended)
- Bash
- Any executable language that returns JSON

Tip: Custom modules should follow the Ansible interface (stdin for input, JSON for output) to integrate smoothly.

Categories of Commonly Used Modules

Modules are grouped by function. Here's a breakdown of the most important and frequently used types:

1. File & Directory Management

Module	Purpose
file	Set permissions, create/remove files and directories
сору	Copy files from control node to target
templ ate	Use Jinja2 templates to create config files

stat Check file attributes

2. Service Management

Module	Purpose
servi	Manage services via
ce	SysV/upstart/systemd
syste md	Modern service management

3. Package Management

Module	Platform	Purpose
apt	Debian/Ubu ntu	Install, remove packages
yum	RHEL/CentO S	Install, remove packages
dnf	Fedora/RHE L	Modern Yum replacement
pip	Python	Install Python packages

1 4. User and Group Management

Module Purpose

user Manage system users

group Manage system groups

5. Networking & Firewalls

6. Cloud and Infrastructure

Provider	Example Modules
AWS	ec2, s3, elb, route53
Azure	azure_rm_*

GCP gcp_compute_instance,

gcp_storage

Docker docker_container,

docker_image

Kubernet k8s, helm

es

7. Utilities and Helpers

Module	Purpose
debug	Print messages in playbooks
set_f act	Define variables dynamically
pause	Pause execution for input/timing
asser t	Validate conditions and fail with message

Modules Work Behind the Scenes

Understanding what happens internally helps you troubleshoot and optimize your playbooks.

X Step-by-Step Process:

- 1. Task is defined in the playbook, using a module.
- 2. Ansible connects to the target host (typically via SSH).
- 3. The module (Python script or executable) is **transferred to the host**.
- 4. The module is **executed on the host**, using the arguments you passed in YAML.
- 5. The module **performs the action**, checks for changes, and returns a **JSON** response.
- 6. Ansible captures the result, updates its state, and shows you output.

Example JSON Output from a Module:

```
{
  "changed": true,
  "msg": "nginx installed",
  "invocation": {
    "module_args": {
```

Why Understanding Modules Matters

- Helps you write better playbooks (more efficient, more readable).
- Makes debugging easier when things go wrong.
- Enables you to **extend** Ansible with custom functionality.
- Builds a foundation for **advanced topics** like roles, collections, and plugins.

Ø Use Case:

You're tasked with setting up a **production-ready Nginx web server** on a group of Ubuntu hosts. The server must:

- Install Nginx
- Serve content from a specific directory
- Have a custom HTML landing page
- Ensure Nginx is enabled and running
- Make sure UFW allows HTTP traffic

Modules Covered:

- apt
- copy
- template
- file

• service • ufw Folder Structure (Recommended) webserver-setup/ — files/ └─ index.html templates/ └─ nginx.conf.j2 playbook.yml flaybook.yml - name: Setup Nginx Web Server hosts: webservers become: yes

vars:

```
web_root: /var/www/html
 nginx_port: 80
tasks:
  - name: Update apt cache
    apt:
      update_cache: yes
      cache_valid_time: 3600
  - name: Install Nginx
    apt:
      name: nginx
      state: present
 - name: Copy custom index.html
   copy:
      src: files/index.html
      dest: "{{ web_root }}/index.html"
```

owner: www-data

group: www-data

mode: '0644'

- name: Deploy custom nginx config

template:

src: templates/nginx.conf.j2

dest: /etc/nginx/sites-available/default

notify: Reload Nginx

- name: Ensure Nginx is started and enabled

service:

name: nginx

state: started

enabled: true

- name: Allow HTTP through UFW

ufw:

rule: allow

```
port: "{{ nginx_port }}"
        proto: tcp
  handlers:
    - name: Reload Nginx
      service:
        name: nginx
        state: reloaded
files/index.html
<!DOCTYPE html>
<html>
<head>
    <title>Welcome to Ansible Web Server</title>
</head>
<body>
    <h1>It works! 4</h1>
</body>
```

```
</html>
```

```
 templates/nginx.conf.j2
```

```
server {
    listen {{ nginx_port }};
    server_name _;
    root {{ web_root }};

    index index.html;

    location / {
        try_files $uri $uri/ =404;
    }
}
```

Why This Task Matters

- Package installation: apt is a go-to module in real-world Ubuntu environments.
- File and template management: copy and template show how to manage real content and configs.
- Service handling: service ensures your services are up and running.
- Firewall configuration: ufw ensures your server is reachable securely.

These are foundational skills in CI/CD pipelines, PaaS infrastructure, and provisioning workflows.

Result:

After running this playbook:

- Your Nginx web server is installed.
- A custom landing page is live.
- A custom config is deployed.
- Firewall allows traffic.
- It's ready for production or containerization.

ℛ Real-World Practical Task #2: Create a New System User with SSH Access for CI/CD Agent

Ø Objective:

Create a dedicated Linux user (ci-agent) for CI/CD purposes. This user should:

- Exist on all target servers
- Have a valid home directory
- Be given a public SSH key for authentication
- Have sudo access (optional, controlled by a variable)

Modules Used:

- user
- authorized_key
- copy
- lineinfile or template (for sudoers)
- group

• stat (optional validation)

Folder Structure

playbook.yml

```
---
- name: Create a CI/CD System User with SSH Access
hosts: all
become: yes

vars:
   ci_user: ci-agent
   ci_home: "/home/{{ ci_user }}"
```

```
ci_ssh_pubkey: "{{ lookup('file', 'files/id_rsa.pub')
}}"
   grant_sudo: true
 tasks:
    - name: Ensure '{{ ci_user }}' group exists
      group:
        name: "{{ ci_user }}"
        state: present
    - name: Create '{{ ci_user }}' user
      user:
        name: "{{ ci_user }}"
        comment: "CI/CD Agent User"
        shell: /bin/bash
        group: "{{ ci_user }}"
        home: "{{ ci_home }}"
        create_home: yes
        state: present
```

```
- name: Create .ssh directory
  file:
    path: "{{ ci_home }}/.ssh"
    state: directory
    owner: "{{ ci_user }}"
    group: "{{ ci_user }}"
    mode: '0700'
- name: Add public SSH key for '{{ ci_user }}'
  authorized_key:
    user: "{{ ci_user }}"
    key: "{{ ci_ssh_pubkey }}"
    state: present
    manage_dir: no # We already created the .ssh dir
- name: Optionally grant sudo access
  lineinfile:
    path: /etc/sudoers.d/{{ ci_user }}
```

```
create: yes
state: present
line: "{{ ci_user }} ALL=(ALL) NOPASSWD:ALL"
mode: '0440'
when: grant_sudo
```

Result:

- You now have a clean, automated, secure way to provision a system user with SSH access.
- This user can be used for CI/CD pipelines, remote scripts, or infrastructure automation.
- With grant_sudo, you can dynamically choose whether the user has privileged access.

Why This Task Is Practical

- Used in every cloud deployment or pipeline bootstrapping step.
- Gives you reusable logic for bootstrapping secure remote users.
- Teaches critical modules (user, authorized_key, lineinfile) used in production and auditing tasks.

Ø Objective:

Use Ansible to install Docker (if not present), deploy a containerized web application, and ensure it auto-starts on system reboot. This could be a simple nginx, httpd, or a real microservice container.

Modules Covered:

- apt (install dependencies)
- docker_container
- docker_image
- service (for Docker daemon)
- uri (optional health check)
 - This task introduces the **Docker modules**, which are heavily used in cloud-native and containerized CI/CD pipelines.

Folder Structure

Anatomy of This Playbook

Section	Purpose
pre_ta	Ensure dependencies are present
tasks	Install Docker, pull image, run container
vars	Define Docker image, port, etc.
handle rs	Restart Docker if config changes

% playbook.yml

- name: Deploy Dockerized Web Application hosts: app_servers become: yes vars: docker_package_state: present docker_service_state: started docker_image_name: nginx docker_image_tag: latest docker_container_name: webapp docker_container_port: 80 host_port: 8080 pre_tasks:

- name: Install required packages

```
apt:
        name:
          - apt-transport-https
          - ca-certificates
          - curl
          - gnupg
          - lsb-release
        state: present
        update_cache: yes
    - name: Add Docker GPG key
      apt_key:
        url: https://download.docker.com/linux/ubuntu/gpg
        state: present
    - name: Add Docker repository
      apt_repository:
        repo: "deb [arch=amd64]
https://download.docker.com/linux/ubuntu {{
ansible_distribution_release }} stable"
```

```
notify: Update APT cache
tasks:
  - name: Install Docker
    apt:
      name: docker-ce
      state: "{{ docker_package_state }}"
    notify: Restart Docker
  - name: Ensure Docker service is running
    service:
      name: docker
      state: "{{ docker_service_state }}"
      enabled: true
  - name: Pull Docker image "{{ docker_image_name }}"
    docker_image:
```

state: present

```
name: "{{ docker_image_name }}"
        tag: "{{ docker_image_tag }}"
        source: pull
    - name: Ensure "{{ docker_container_name }}" container
is running
      docker_container:
        name: "{{ docker_container_name }}"
        image: "{{ docker_image_name }}:{{ docker_image_tag
}}"
        state: started
        restart_policy: always
        published_ports:
          - "{{ host_port }}:{{ docker_container_port }}"
    - name: (Optional) Health check for the web app
      uri:
        url: "http://localhost:{{ host_port }}"
        return_content: yes
        status_code: 200
```

```
register: web_health
    retries: 3
    delay: 5
    until: web_health.status == 200
handlers:
  - name: Restart Docker
    service:
      name: docker
      state: restarted
  - name: Update APT cache
    apt:
      update_cache: yes
```

What This Playbook Achieves

• Installs and configures Docker on any Ubuntu host

- Deploys a Dockerized nginx container exposed on port 8080
- Ensures it auto-restarts on reboot or crash
- Optionally performs a live health check on the container

Why This Task Is DevOps Gold

- Recreates an end-to-end delivery of an app in a cloud-native way
- Embraces infrastructure-as-code, containerization, and self-healing
- Builds foundational knowledge to integrate with CI/CD, Kubernetes, ECS,
 EKS, and more

Ø Objective:

Set up **HAProxy** on a designated load balancer server to distribute traffic to a backend pool of web servers. This setup is essential in high-availability, scalable infrastructure environments.

Modules Covered:

- apt Install HAProxy
- template Deploy HAProxy config dynamically
- service Manage HAProxy
- lineinfile For optional sysctl/network tuning
- copy Optional: deploy SSL certs or error pages

Use Case Summary:

- One **HAProxy node** acts as a load balancer
- Backend web servers (e.g., nginx) are assumed to be up and running
- Load balancer listens on port 80
- Balances HTTP traffic across multiple web servers

Folder Structure

```
haproxy-setup/

— playbook.yml

— templates/

— haproxy.cfg.j2
```

playbook.yml

- name: Configure HAProxy Load Balancer

hosts: loadbalancer

```
become: yes
vars:
  haproxy_cfg_path: /etc/haproxy/haproxy.cfg
  frontend_port: 80
 backend_servers:
    - name: web1
      address: 192.168.56.101
      port: 80
    - name: web2
      address: 192.168.56.102
      port: 80
```

tasks:

- name: Install HAProxy
apt:
name: haproxy

state: present

```
- name: Deploy HAProxy configuration
    template:
      src: templates/haproxy.cfg.j2
      dest: "{{ haproxy_cfg_path }}"
      mode: '0644'
    notify: Restart HAProxy
  - name: Ensure HAProxy service is running and enabled
    service:
      name: haproxy
      state: started
      enabled: true
handlers:
  - name: Restart HAProxy
    service:
```

update_cache: yes

name: haproxy

state: restarted

templates/haproxy.cfg.j2

```
global
  log /dev/log local0
  log /dev/log local1 notice
  daemon
  maxconn 256
```

defaults

```
log global
mode http

option httplog

option dontlognull
retries 3

timeout connect 5000ms

timeout client 50000ms
```

```
timeout server 50000ms
```

```
frontend http_front
  bind *:{{ frontend_port }}
  default_backend http_back

backend http_back
{% for srv in backend_servers %}
  server {{ srv.name }} {{ srv.address }}:{{ srv.port }}
check
{% endfor %}
```

Outcome:

- A ready-to-use HAProxy load balancer
- Distributes incoming HTTP traffic between backend nodes
- Automatically reconfigurable via template updates
- Ensures HAProxy starts on boot and restarts when the config changes

Real-Life Application:

- Used in production for high availability and failover
- Common in microservices and monolith-to-microservice transition environments
- Great foundation for introducing TLS termination, health checks, or blue-green deployments

Real-World Practical Task #5: Automated User Management with Ansible

Ø Objective:

Use Ansible to:

- Create multiple users across multiple servers
- Add SSH public keys for secure access
- Assign (or revoke) sudo privileges
- Remove unwanted users safely

This is one of the most **frequently automated tasks** in infrastructure management, especially in **team onboarding/offboarding** or **multi-node fleet configuration**.

Modules Covered:

- user
- authorized_key
- group

- lineinfile
- copy

Folder Structure

group_vars/all.yml

users_to_create:

```
- username: john
shell: /bin/bash
ssh_key: "{{ lookup('file', 'files/john.pub') }}"
sudo: true
```

```
- username: alice
    shell: /bin/zsh
    ssh_key: "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAABAQD..." #
paste her key directly
    sudo: false
users_to_remove:
  - mark
  - tempuser
% playbook.yml
- name: Manage system users
  hosts: all
  become: yes
  tasks:
    - name: "Create user '{{ item.username }}'"
```

```
user:
    name: "{{ item.username }}"
    shell: "{{ item.shell | default('/bin/bash') }}"
    state: present
    create_home: yes
  loop: "{{ users_to_create }}"
  tags: [create, users]
- name: "Add SSH key for '{{ item.username }}'"
  authorized_key:
    user: "{{ item.username }}"
    key: "{{ item.ssh_key }}"
    state: present
  loop: "{{ users_to_create }}"
  when: item.ssh_key is defined
  tags: [ssh, users]
- name: "Add '{{ item.username }}' to sudoers"
  lineinfile:
```

```
path: /etc/sudoers.d/{{ item.username }}
    create: yes
    mode: '0440'
    line: "{{ item.username }} ALL=(ALL) NOPASSWD:ALL"
 when: item.sudo | default(false)
  loop: "{{ users_to_create }}"
  tags: [sudo, users]
- name: "Remove user '{{ item }}'"
  user:
    name: "{{ item }}"
    state: absent
    remove: yes
  loop: "{{ users_to_remove }}"
  tags: [delete, users]
```

What This Playbook Does

Action

- Creates specified users with custom shells
- Installs their SSH keys for secure login
- Grants sudo access when requested
- Removes unwanted users cleanly (with home dirs)

Why This Task is DevOps-Friendly

- **Idempotent** rerunning it won't break existing users
- Modular add/remove users via group_vars/all.yml only
- Secure handles sudo access through dedicated files
- **Extensible** easy to integrate into CI/CD pipelines or cloud-init bootstraps

® Real time Task 7:

Use Ansible to:

- Perform safe and controlled unattended OS patching
- Detect and notify if a reboot is required
- Allow scheduling during maintenance windows

Why This Matters in Real-World Environments

- Security compliance (e.g., CIS, PCI-DSS, HIPAA)
- Reduce human error during updates
- Avoid unplanned downtime with **controlled reboots**
- Schedule updates in low-traffic windows

Task Breakdown

- 1. Run safe updates (apt or yum)
- 2. Notify/report if a reboot is required
- 3. Optionally reboot (manually or automatically)
- 4. Ensure idempotency and easy logging

Modules Used

- apt or yum Perform upgrades
- shell or stat Check for reboot-required files
- reboot Handle reboots (optionally)
- ansible_facts Access system info
- cron or at (optional) Schedule updates in windows

Example Playbook (Ubuntu / Debian)

- name: Patch Management - Safe OS Updates hosts: all become: yes gather_facts: yes vars: auto_reboot: false maintenance_window: "02:00" tasks: - name: Update package cache apt: update_cache: yes cache_valid_time: 3600

```
- name: Upgrade all packages safely
 apt:
   upgrade: safe
- name: Check if reboot is required (Debian/Ubuntu)
  stat:
    path: /var/run/reboot-required
  register: reboot_required
- name: Notify if reboot is needed
 debug:
   msg: "ॡ Reboot is required!"
 when: reboot_required.stat.exists
- name: Reboot the machine (optional)
  reboot:
   msg: "Reboot initiated by Ansible after patching"
   pre_reboot_delay: 30
    reboot_timeout: 600
```

```
when: reboot_required.stat.exists and auto_reboot
    - name: Log result to local file (optional)
      local_action:
        module: copy
        content: |
          Host: {{ inventory_hostname }}
          Reboot Needed: {{ reboot_required.stat.exists }}
          Updated At: {{ ansible_date_time.date }} {{
ansible_date_time.time }}
        dest: "./patch_logs/{{ inventory_hostname }}.log"
RHEL/CentOS Variant (yum)
Replace the apt tasks with:
- name: Upgrade all packages on RHEL/CentOS
  yum:
    name: "*"
    state: latest
```

And check for reboot:

```
- name: Check if reboot hint exists (RHEL)
stat:
   path: /var/run/reboot-required
register: reboot_required
```

Or use:

```
- name: Check if last kernel version is in use
shell: needs-restarting -r
register: reboot_hint
failed_when: false
```

Optional: Patch Only in Maintenance Windows

Use cron or at to delay patch tasks (or run playbook only if ansible_date_time.hour == "02").



Description Optional: Reporting

Use mail, slack, or uri module to send patch results to your team. Example:

```
- name: Send notification to Slack
  uri:
    url: https://hooks.slack.com/services/XXX
    method: POST
    body: '{"text": "Reboot required on {{
inventory_hostname }}"}'
    headers:
      Content-Type: "application/json"
 when: reboot_required.stat.exists
```

Best Practices

- Tag playbooks with patch for easy scheduling: ansible-playbook patch.yml --tags "patch"
- Use inventory groups like [prod], [staging] for phased rollouts
- Store logs centrally (e.g., in S3, log server)

© Goal of This Task

This Ansible task automates the **entire OS patching process** in a clean, controlled, and auditable way. It's especially useful in environments where:

- Systems must stay secure and up-to-date
- Reboots need to be tracked and minimized
- Maintenance windows must be respected
- Patching should happen without manual SSH login

Step-by-Step Explanation

1. Update the Package Cache

- name: Update package cache
apt:
 update_cache: yes
 cache_valid_time: 3600

Q Why?

Before upgrading, we make sure the system knows about the latest package versions. This avoids stale package data, which could cause failed upgrades.

2. Perform a Safe Upgrade

- name: Upgrade all packages safely
 apt:
 upgrade: safe
- What is "safe"?
 - Only upgrades **already installed** packages
 - Doesn't remove packages or install new dependencies
 - Less risky than full-upgrade or dist-upgrade
- Real-world Benefit: This keeps systems secure without unexpected service disruptions.

3. Check If a Reboot Is Required

- name: Check if reboot is required (Debian/Ubuntu)
stat:
 path: /var/run/reboot-required
register: reboot_required

Why check this?

Many Linux systems drop a file here after kernel updates or critical patches — it's a **standard signal** that a reboot is needed for changes to take effect.

4. Notify If Reboot Is Needed

- name: Notify if reboot is needed
 debug:
 msg: " Reboot is required!"
 when: reboot_required.stat.exists

____ This will log or print a message if a reboot is necessary. In real usage, you might integrate this with Slack, email, or another alerting system.

5. (Optional) Reboot the Server

```
- name: Reboot the machine (optional)
  reboot:
    msg: "Reboot initiated by Ansible after patching"
    pre_reboot_delay: 30
    reboot_timeout: 600
when: reboot_required.stat.exists and auto_reboot
```

Controlled reboot

- Only happens if a reboot is required and you've explicitly allowed it with auto_reboot: true
- Waits 30 seconds, gives system time to close services
- Waits up to 600 seconds (10 minutes) for server to come back
- Great for **automated but safe** reboot workflows

6. (Optional) Log the Result

```
- name: Log result to local file (optional)
local_action:
   module: copy
...
```

- Saves patch and reboot information in a file on the control node. Useful for:
 - Audits
 - Troubleshooting
 - Compliance reporting

A Optional Enhancements

Time-Windowed Patching

This task can easily be adapted to **only run during specific hours** (e.g., 2–3 AM) using:

• cron to schedule the playbook

when: conditions based on system time (e.g.,
ansible_date_time.hour == '02')

Notifications

Add uri or mail modules to send alerts if:

- A reboot is required
- Patching fails
- Patching completes successfully

Best Practices for Ansible Playbooks &

Modules

Ansible is powerful, but with great power comes great responsibility. Following best practices ensures that your automation is:

- Readable
- Reusable
- Reliable
- Idempotent
- Easy to scale & debug

1. Keep Playbooks Readable and Modular

Use roles to organize logic

Break playbooks into reusable **roles** for tasks like nginx, users, mysql, etc.

```
site.yml
roles/
  ├─ nginx/
     users/
```

```
└─ mysql/
```

Separate configuration (vars) from logic (tasks)

Bad:

```
- name: Install nginx
  apt: name=nginx state=present update_cache=yes
Good:
vars:
  nginx_package: nginx
tasks:
  - name: Install nginx
    apt:
      name: "{{ nginx_package }}"
      state: present
      update_cache: yes
```

🧪 2. Write Idempotent Tasks

Ansible is declarative. Every task should be safe to run multiple times without causing change.

Good (idempotent):

- name: Add user

user:

name: deploy

state: present

X Bad (not idempotent):

- name: Create user with shell

shell: useradd deploy

3. Use Ansible Modules Over Shell Commands

Always prefer Ansible modules — they're idempotent, easier to read, and error-aware.

Good:

- name: Install packages

apt:

```
name: "{{ item }}"
state: present
loop:
  - git
  - curl
```

X Bad:

```
- name: Install packages (bad)
shell: apt install -y git curl
```

4. Variable Management

- Use defaults/main.yml in roles for sane defaults
- Use **group_vars/host_vars** for targeted overrides
- Avoid hardcoded values
- Use vars_prompt for sensitive or interactive input when necessary

5. Inventory & Environment Structure

Maintain clear environment-based inventories:

```
inventories/
production/
hosts.yml
group_vars/
staging/
hosts.yml
group_vars/
```

Use --inventory to target environment.

6. Use Loops and Conditionals Wisely

Example: Looping with when

```
- name: Install only if on Debian
apt:
    name: nginx
    state: present
```

```
when: ansible_os_family == "Debian"
```

🔒 7. Handle Secrets Securely

```
Use Ansible Vault for passwords, tokens, SSH keys:
ansible-vault encrypt secrets.yml
```

```
- name: Use encrypted secret
  debug:
    msg: "My password is {{ vault_db_pass }}"
```

8. Test Your Playbooks

- Use --check for dry runs
- Use --diff to show changes
- Use ansible-lint to catch bad patterns
- Test playbooks in Vagrant, Docker, or EC2 test environments

9. Use Meaningful Names and Tags

- Give every task a clear name
- Use tags to run parts of a playbook:

```
- name: Install nginx
apt:
   name: nginx
   state: present
tags: nginx
```

Run: ansible-playbook site.yml --tags nginx

📑 10. Documentation Inside Your Playbooks

- Use name: fields everywhere
- Use comments only when the logic isn't obvious
- Document README.md per role

11. Limit Host Impact and Use Serial

• Use serial for rolling changes:

- hosts: webservers

serial: 2

This prevents downtime across all servers during patching or deployments.

® Bonus Tips:

1. become: true Only Where Needed

☑ Best Practice:

Use privilege escalation (become: true) only on tasks that absolutely require root access.

- name: Install NGINX

become: true

apt:

name: nginx

state: present

M Why it matters:

- Improves security posture
- Prevents accidental destructive changes (especially when variables are misconfigured)
- Aligns with the principle of **least privilege**

• Apply become: true at the task level (not globally), unless all tasks in the play require it.

2. Use check_mode: in CI

W Best Practice:

Run your playbooks in --check mode during testing pipelines.

ansible-playbook site.yml --check

You can also selectively disable check mode for specific tasks:

- name: Restart service (skip in check mode)

service:

name: nginx

state: restarted

check_mode: no

6 Why it matters:

- Validates syntax and structure without making real changes
- Ensures **idempotency** (Ansible shows what *would* change)

• Great fit for CI/CD pipelines, especially when automating infrastructure

3. Use block, rescue, always for Task Control

This allows **error handling**, much like try/except/finally in code.

```
- name: Install with error handling
  block:
    - name: Try installing custom package
      apt:
        deb: /tmp/custom-package.deb
  rescue:
    - name: Fallback to APT install
      apt:
        name: fallback-package
        state: present
  always:
    - name: Clean temp files
      file:
        path: /tmp/custom-package.deb
        state: absent
```

Why it matters:

- Gracefully handles errors without breaking the entire playbook
- Ensures cleanup always runs
- Allows you to attempt multiple strategies (like retrying a failed download)

4. Reuse with include_tasks and import_tasks

Use task inclusion to break large playbooks into logical units.

```
# In main.yml
```

- name: Install web stack

import_tasks: web.yml

- name: Set up users

include_tasks: users.yml

- import_tasks: Static include (resolved at parse time)
- include_tasks: Dynamic include (resolved at runtime, can use variables)

Why it's powerful:

- Encourages DRY (Don't Repeat Yourself) principles
- Makes debugging easier with smaller, modular files
- Promotes task reusability across multiple playbooks
 - Use include_tasks when task names depend on conditionals or variables.

5. Store Reusable Roles in Ansible Galaxy or Private GitRepos

- Package reusable logic into roles
- Host them in:
 - Ansible Galaxy (ansible-galaxy publish)
 - Private Git repos (requirements.yml)

```
# requirements.yml
```

- src: git+https://github.com/my-org/ansible-role-nginx.git

name: nginx

version: v1.0.3

Install with:

ansible-galaxy install -r requirements.yml

Why this scales:

- Makes roles versioned and maintainable
- Promotes **collaboration** in DevOps teams
- Easy to **pin, update, or reuse** in other environments/projects

X Troubleshooting Tips in Ansible

Whether you're working on a single-node test or orchestrating across dozens of production hosts, issues happen. The key is knowing **how to trace, diagnose, and fix** them quickly.

1. Use −v, −vv, or −vvv for Verbose Output

Ansible gives **minimal output by default**, which is great for clean logs but not for debugging.

- -v: shows task-level info (like changes made)
- -vv: shows module arguments & connection details
- -vvv: shows SSH-level debug, raw output

bash

CopyEdit

ansible-playbook site.yml -vvv

Start with -vv for most issues.

2. Check the Inventory File

Common issue: hosts not found, or groups incorrectly targeted.

- Ensure your inventory.ini or hosts.yaml is properly formatted
- Confirm hostnames/IPs are reachable
- Validate group names match the playbook

```
ansible-inventory --list -y
```

 \bigvee Pro Tip: Use --limit to narrow scope while testing: ansible-playbook site.yml --limit web1.example.com

3. Use --check Mode for Dry Runs

The --check flag simulates the playbook without making changes. It helps validate logic, task flow, and conditional behavior.

ansible-playbook nginx.yml --check

Limitations:

Not all modules support check mode

• Doesn't show file/template content differences

4. Clean Up Facts & Cache

Issues from outdated facts can confuse playbooks.

✓ Clear fact cache:

ansible all -m setup --flush-cache

Or disable fact caching in your config or playbook:

gather_facts: no

Tip: Only gather facts if you need them.

5. Use ansible-lint to Catch Common Errors

ansible-lint identifies bad practices, typos, or dangerous usage.
ansible-lint playbook.yml

- Flags deprecated modules
- Warns on shell usage when a module exists

• Ensures consistent syntax and style



Use the debug module:

```
- debug:
    var: some_variable

Or print a message:

- debug:
    msg: "Deploying to {{ inventory_hostname }}"
```

Use when: to narrow down tasks and confirm logic is working.

🧠 7. Trace with Tags and --start-at-task

- Use --tags to run a subset of tasks
- Use --start-at-task to resume long playbooks without re-running everything

ansible-playbook site.yml --start-at-task="Install NGINX"

8. Check Permissions & become Issues

Common problems:

- Task fails with Permission denied
- Missing become: true or become_user

Example:

- name: Restart service

service:

name: nginx

state: restarted

become: true

Also check:

- User in inventory (ansible_user)
- SSH key permissions
- sudo access for remote user

9. Common Error Examples & Fixes

Error Message	Likely Cause	Fix
UNREACHABLE!	SSH/auth failure	Check IP, user, key, port
<pre>FAILED! => no attribute</pre>	Bad variable reference	Check spelling, default/fallback
File not found	Relative path problem	<pre>Use full path or {{ role_path }}/files/filename</pre>
Missing required parameters	Incomplete module args	Refer to module docs
Permission denied	No root access	Use become: true, verify sudoers

📚 10. Use ansible-playbook --syntax-check

Catches YAML syntax issues or malformed playbooks:

ansible-playbook site.yml --syntax-check

11. Inspect Host-Specific Facts

Sometimes a task behaves differently per host. Use the setup module to inspect:

ansible web1 -m setup | less

Search for specific facts:

ansible web1 -m setup -a 'filter=ansible_distribution'

12. Check Ansible Configuration

Run this to see current config, overrides, and paths:

ansible-config dump --only-changed

Useful to debug:

Inventory plugin issues

- Callback plugin behavior
- Fact caching settings
- SSH control settings

© Final Thoughts & What's Next

Congratulations—you've reached the end of Automate Like a Pro: The DevOps Guide to Ansible Playbooks & Modules!

Over the course of this doc, you've:

- Gained deep understanding of core Ansible concepts and structure
- Built practical, real-world automation tasks that DevOps engineers use daily
- Learned best practices to make your playbooks clean, secure, and reusable
- Added debugging and troubleshooting techniques that empower you in production

** What to Look Forward To

This guide is only one part of a much larger journey—here's what's coming soon:

- Advanced Ansible Techniques: Roles, Collections, Plugins
- CI/CD Integrations: GitLab, Jenkins, GitHub Actions
- Container & Orchestration Workflows: Docker Compose, Kubernetes
- Cloud-Native Patterns: AWS/Azure/GCP provisioning with infrastructure as code
- Security & Compliance Automation: Vault, hardening pipelines, audit-ready systems

Stay tuned—there's more powerful, hands-on Ansible wisdom coming your way very soon!

Fenil Gajjar :

"Knowledge grows only when shared. Let's walk together, grow together, and lead by example."